

DevOps Project Final Report

Project: Book Library REST API with Full DevOps Implementation

Author: Ayhem Boukari

1. Project Overview

This project demonstrates a comprehensive DevOps implementation for a Python Flask REST API. The Book Library API provides CRUD operations for managing books with full observability, containerization, Kubernetes orchestration, and automated CI/CD pipelines.

Key Achievements:

- Built a production-ready REST API in under 150 lines of code
- Implemented complete observability stack (metrics, logs, tracing)
- Created multi-stage Docker build with optimized image size
- Deployed to Kubernetes with health checks and auto-scaling capabilities
- Automated testing, security scanning, and deployment through GitHub Actions

2. Architecture

System Components

The architecture follows cloud-native principles with clear separation of concerns:

Component	Technology	Purpose
API Server	Flask + Gunicorn	RESTful endpoints with WSGI server
Metrics	Prometheus Client	Expose application metrics
Monitoring	Prometheus + Grafana	Collect and visualize metrics
Container Runtime	Docker	Application containerization
Orchestration	Kubernetes	Container orchestration and scaling
CI/CD	GitHub Actions	Automated testing and deployment

Design Decisions

1. **Flask Framework:** Chosen for simplicity and low overhead, perfect for microservices
2. **Gunicorn WSGI:** Production-grade server with worker process management
3. **In-Memory Storage:** Simplified data layer for demonstration (easily replaceable with database)
4. **Multi-Stage Docker:** Reduces image size from ~900MB to <150MB
5. **Prometheus Metrics:** Industry-standard monitoring compatible with Kubernetes

3. CI/CD Pipeline

Pipeline Architecture

```
Code Push → Lint → Test → Build → Security Scan → Deploy → Verify
```

Pipeline Stages

Stage	Tools	Purpose
Lint	flake8, black	Code quality and style enforcement
Test	pytest, coverage	Unit tests with 90%+ coverage
Build	Docker Buildx	Multi-platform image building
Security	Bandit, Semgrep, Trivy, ZAP	SAST, DAST, and container scanning
Deploy	kubectl	Kubernetes deployment with rollout verification

Automation Benefits

- **Consistency:** Every change goes through identical validation
- **Speed:** Parallel job execution reduces pipeline time to ~5 minutes
- **Safety:** Automated rollback on deployment failure
- **Visibility:** Artifact uploads for all security reports

4. Containerization

Docker Strategy

The Dockerfile uses a multi-stage build pattern:

1. **Builder Stage:** Installs dependencies in a virtual environment
2. **Production Stage:** Copies only necessary files with non-root user

Optimizations

Optimization	Impact
Multi-stage build	Image size: 890MB → 145MB
Non-root user	Enhanced security posture
Python slim base	Minimal attack surface
.dockerignore	Faster builds, smaller context
HEALTHCHECK	Container self-monitoring

Docker Compose Setup

Local development includes:

- Flask API with hot-reload capability
 - Prometheus for metrics collection
 - Grafana with pre-configured dashboards
-

5. Kubernetes Deployment

Manifests Created

Manifest	Purpose
<code>namespace.yaml</code>	Logical isolation for all resources
<code>configmap.yaml</code>	External configuration management
<code>deployment.yaml</code>	Pod specification with 2 replicas
<code>service.yaml</code>	Internal load balancing (ClusterIP)
<code>ingress.yaml</code>	External access with host-based routing

Deployment Features

- **High Availability:** 2 replicas with anti-affinity rules
- **Rolling Updates:** Zero-downtime deployments with surge/unavailable controls
- **Resource Management:** CPU/memory requests and limits defined
- **Health Probes:** Liveness (/health) and readiness (/ready) checks
- **Security:** Non-root container, read-only filesystem, dropped capabilities

Scaling Considerations

```
resources:
  requests:
    cpu: "100m"
    memory: "128Mi"
  limits:
    cpu: "200m"
    memory: "256Mi"
```

With these limits, each pod can handle ~100 requests/second. Horizontal Pod Autoscaler (HPA) can be added for automatic scaling.

6. Observability

Metrics Implementation

Three key metrics exposed via `/metrics`:

Metric	Type	Labels	Use Case

Metric	Type	Labels	Use Case
http_requests_total	Counter	method, endpoint, status	Request volume and errors
http_request_duration_seconds	Histogram	-	Latency percentiles (P50, P95, P99)
books_total	Gauge	-	Current database state

Logging Strategy

Structured JSON logging with consistent fields:

- **timestamp**: ISO 8601 format for time-series analysis
- **level**: DEBUG/INFO/WARNING/ERROR for filtering
- **request_id**: UUID for request correlation
- **method/path**: Request identification
- **status_code**: Response classification
- **response_time_ms**: Performance tracking

Tracing Approach

Every request receives a unique **X-Request-ID** (UUID v4):

1. Generated at request start
2. Attached to all log entries
3. Returned in response header
4. Enables end-to-end request tracking across services

Grafana Dashboard

Pre-built dashboard includes:

- Real-time request rate by endpoint
- Latency percentile graphs
- Status code distribution
- Books count over time

7. Security

SAST Implementation

Tool	Configuration	Findings
Bandit	Default rules for Python	No high-severity issues
Semgrep	p/python, p/flask rulesets	Flask-specific checks passed

DAST Implementation

OWASP ZAP baseline scan configured to:

- Test all API endpoints
- Check for common vulnerabilities (XSS, injection, headers)
- Generate HTML/JSON reports as artifacts

Dependency Scanning

Tool	Purpose
pip-audit	PyPI vulnerability database
safety	Safety DB for Python packages
Trivy	Container image CVE scanning

Security Improvements Made

1. Non-root container user (UID 1000)
2. Read-only root filesystem in Kubernetes
3. All capabilities dropped
4. Minimal base image (python:slim)
5. Weekly automated security scans

8. Challenges & Lessons Learned

Challenge 1: Keeping Code Under 150 Lines

Solution: Used decorators for cross-cutting concerns (metrics, logging), leveraged Flask's built-in features, and kept the data model simple.

Challenge 2: Docker Image Size

Solution: Multi-stage builds, .dockerignore optimization, and using slim Python base reduced size by 85%.

Challenge 3: Structured Logging Without Library Bloat

Solution: python-json-logger provides JSON formatting with minimal overhead (~10KB).

Challenge 4: Kubernetes Health Probes

Solution: Separated liveness (/health) and readiness (/ready) endpoints with appropriate timeouts to prevent restart loops during slow starts.

Key Takeaways

1. **Infrastructure as Code:** All configurations are version-controlled and reproducible
2. **Shift-Left Security:** Catching vulnerabilities in CI prevents production issues
3. **Observability First:** Built-in from day one, not retrofitted
4. **Automation Everywhere:** Manual processes are eliminated where possible

9. Future Improvements

Short-Term

- Add Horizontal Pod Autoscaler (HPA) for automatic scaling
- Implement database persistence (PostgreSQL)
- Add rate limiting and API authentication

Medium-Term

- Distributed tracing with OpenTelemetry/Jaeger
- Service mesh integration (Istio/Linkerd)
- GitOps deployment with ArgoCD

Long-Term

- Multi-cluster deployment
 - Chaos engineering with Chaos Monkey
 - Cost optimization dashboards
-

Appendix: Quick Reference

Make Commands

```
make help          # Show all commands
make install-dev  # Install dependencies
make test          # Run tests
make lint          # Check code quality
make docker-run    # Start local environment
make k8s-deploy    # Deploy to Kubernetes
make security-scan # Run security scans
```

URLs

Service	URL
API	http://localhost:5000
Prometheus	http://localhost:9090
Grafana	http://localhost:3000

End of Report